



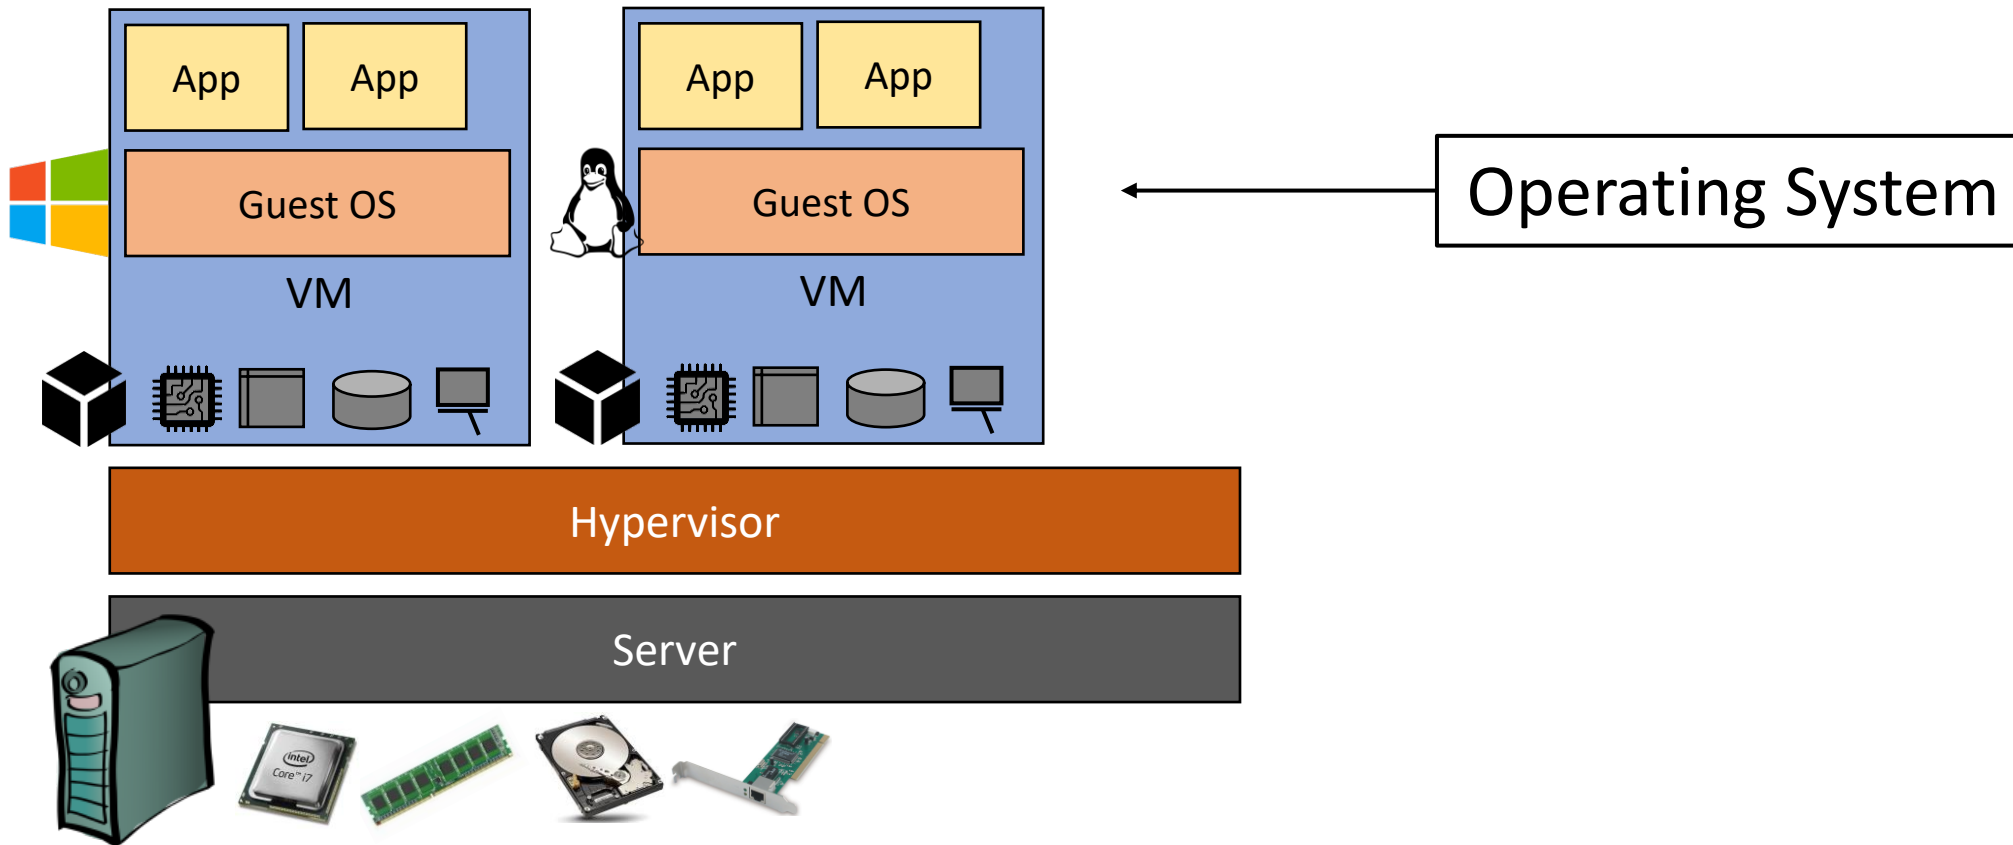
CompSci 401: Cloud Computing

# Overhead of Virtual Machines

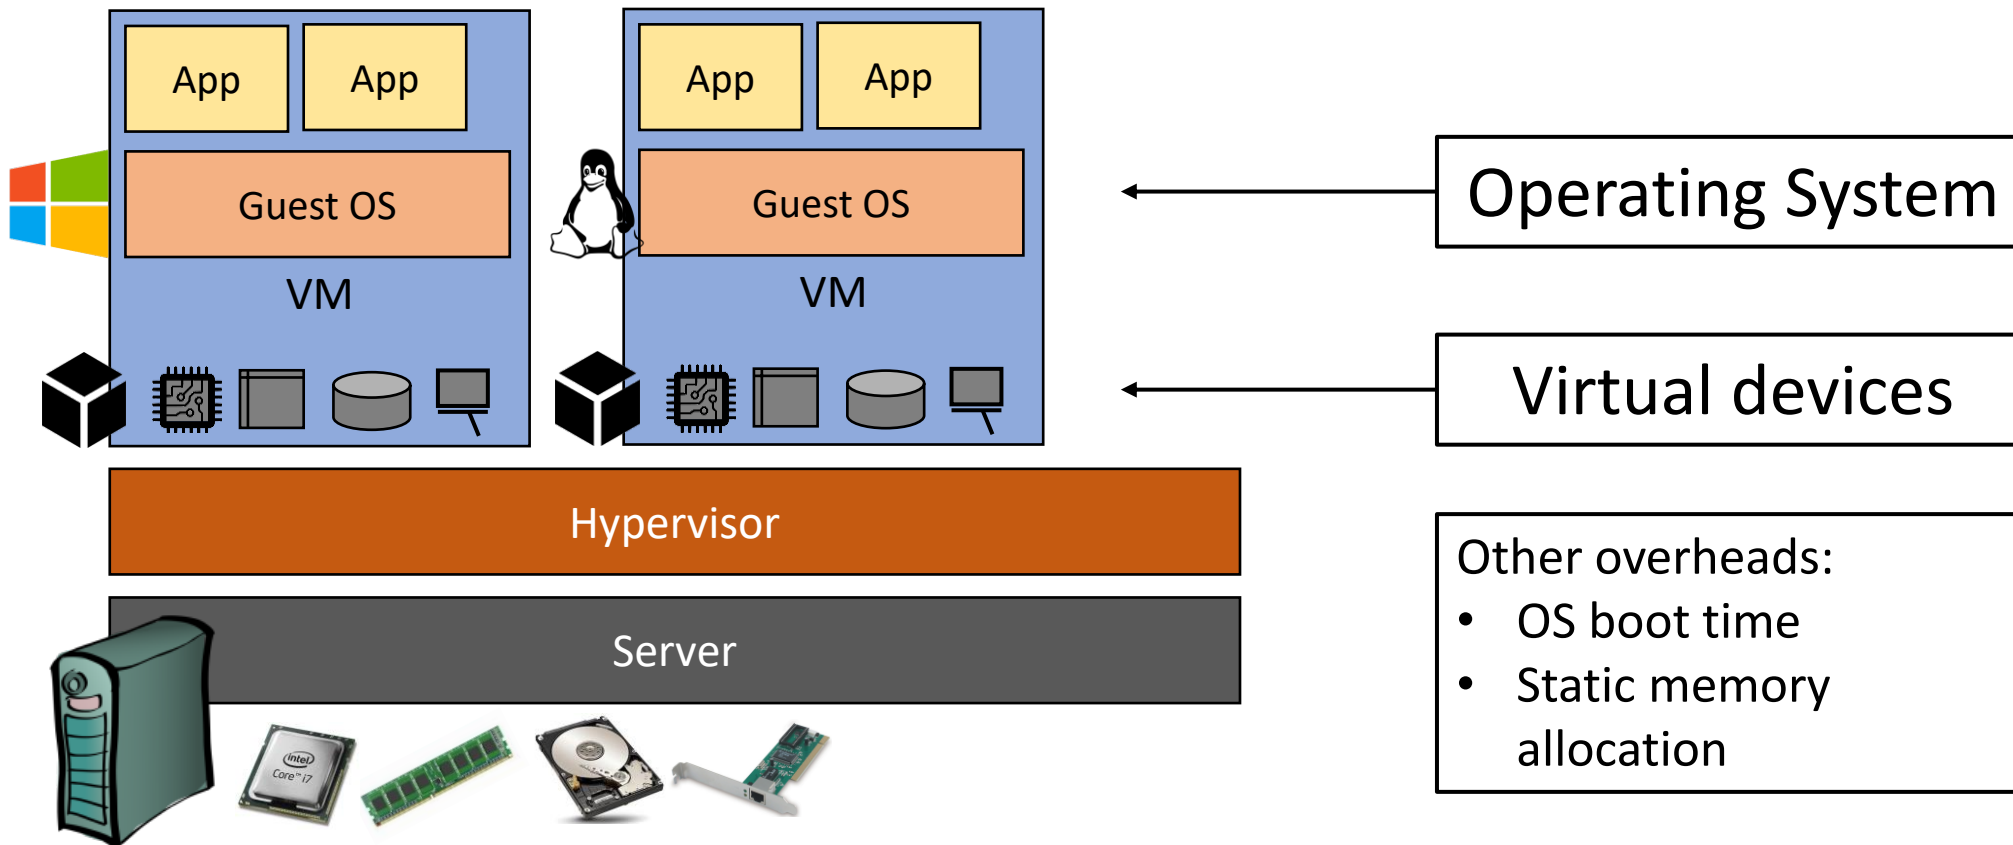
Prof. Ítalo Cunha



# VMs provide complete flexibility



# VMs provide complete flexibility, at a cost



# Virtual machines and autoscaling

- VM overhead is amortized when applications run permanently
  - Pay overhead for booting the operating system once
- VM overhead is amortized when multiple applications run in the OS
  - Pay overhead of the operating system once

# Virtual machines and autoscaling

- VM overhead is amortized when applications run permanently
  - Pay overhead for booting the operating system once
- VM overhead is amortized when multiple applications run in the OS
  - Pay overhead of the operating system once
- Autoscaling means applications do not run permanently
- Most deployments have a single application

# Virtual machines and autoscaling

- VM overhead is amortized when applications run permanently
  - Pay overhead for booting the operating system once
- VM overhead is amortized when multiple applications run in the OS
  - Pay overhead of the operating system once
- Autoscaling means applications do not run permanently
- Most deployments have a single application

**Virtual machines are not a great match for highly dynamic cloud applications**






# CompSci 401: Cloud Computing

# Process Management

Prof. Ítalo Cunha



# Operating systems manage applications






# Operating systems manage applications

## Applications in an operating system

- Start and terminate quickly
- More direct, efficient access to hardware
- Share memory effectively
- Can scale running instances on multiple servers




# Operating systems manage applications

## Applications in an operating system

- Start and terminate quickly
- More direct, efficient access to hardware
- Share memory effectively
- Can scale running instances on multiple servers

OSes do not enforce complete isolation



# Operating systems manage applications


## Applications in an operating system

- Start and terminate quickly
- More direct, efficient access to hardware
- Share memory effectively
- Can scale running instances on multiple servers

## OSes do not enforce complete isolation

Isolation facilities in operating systems:

- Virtual address spaces isolate memory
- User IDs and filesystem permissions isolate data




# Operating systems manage applications

## Applications in an operating system

- Start and terminate quickly
- More direct, efficient access to hardware
- Share memory effectively
- Can scale running instances on multiple servers

## OSes do not enforce complete isolation

- Shared access to the network (IP address, ports)
- Shared filesystem
- Shared metadata
  - What processes are running






# Operating systems manage applications

## Applications in an operating system

- Start and terminate quickly
- More direct, efficient access to hardware
- Share memory effectively
- Can scale running instances on multiple servers

## OSes do not enforce complete isolation

- Shared access to the network (IP address, ports)
- Shared filesystem
- Shared metadata
  - What processes are running
- Free competition to access hardware (memory, CPU)




# Operating systems manage applications

## Applications in an operating system

- Start and terminate quickly
- More direct, efficient access to hardware
- Share memory effectively
- Can scale running instances on multiple servers

## OSes do not enforce complete isolation

- Shared access to the network (IP address, ports)
- Shared filesystem
- Shared metadata
  - What processes are running
- Free competition to access hardware (memory, CPU)



**Sharing leaks information from one application to another**



# CompSci 401: Cloud Computing


# Containers

Prof. Ítalo Cunha



# Shared resources within operating systems


- Operating systems manages resources that are shared across apps
- Protection and isolation mechanisms are not thorough






# Containers have dedicated resources

- Containers do not share any resources with the host OS




# Containers have dedicated resources

- Containers do not share any resources with the host OS
- Linux namespaces:
  - Mount (filesystem)
  - Processes
  - Network
  - Interprocess communication
  - Users
  - Name resolution



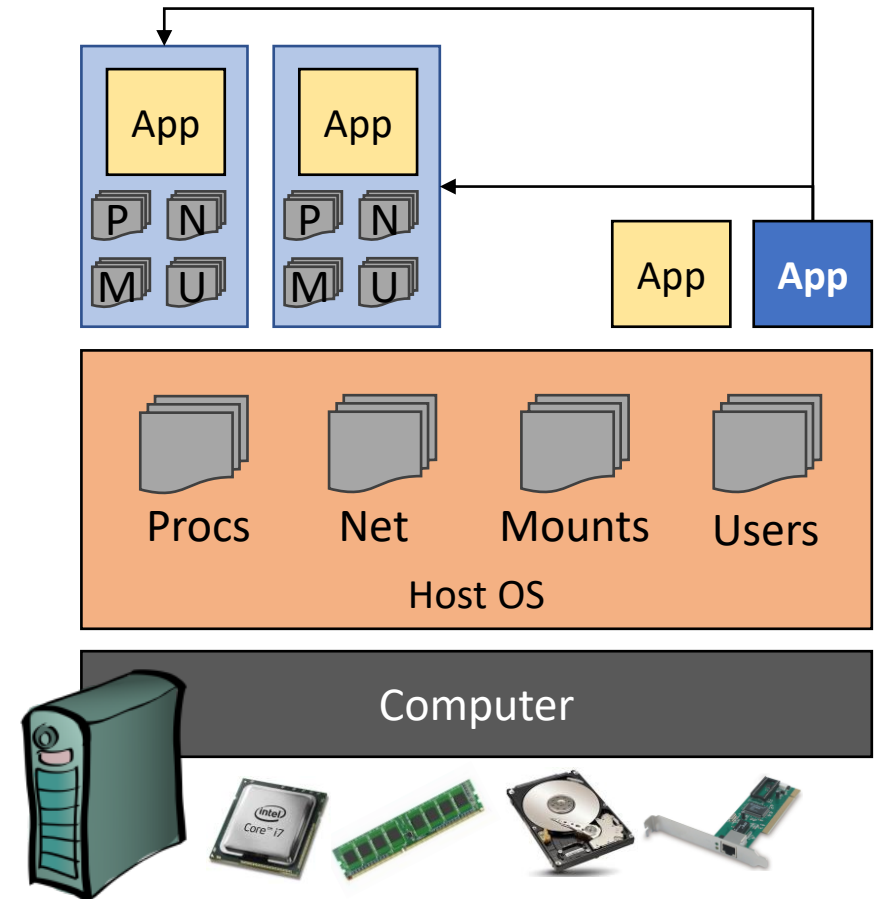
# Containers have dedicated resources

- Containers do not share any resources with the host OS
- Linux namespaces:
  - Mount (filesystem)
  - Processes
  - Network
  - Interprocess communication
  - Users
  - Name resolution
- Resource allocation



# Controlled sharing of resources

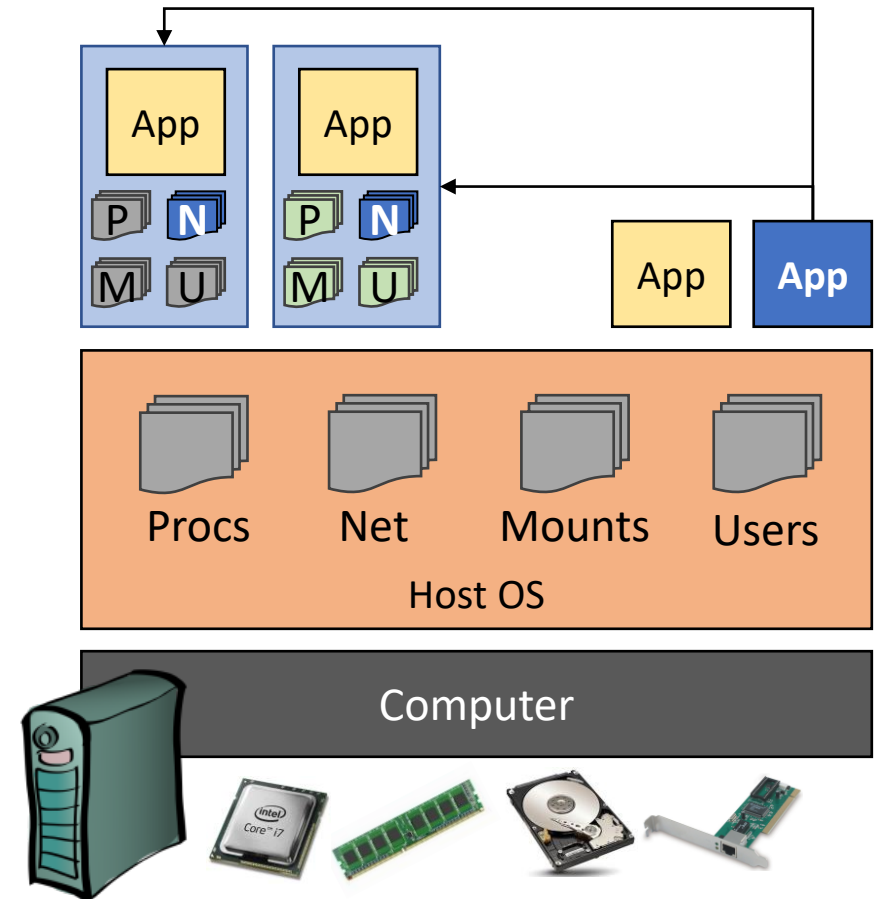
- **Privileged** applications on the host can manage and inspect containers
  - Usually limited to management software like Docker or LXC







# Controlled sharing of resources

- **Privileged** applications on the host can manage and inspect containers
  - Usually limited to management software like Docker or LXC
- Containers can be configured to **share namespaces**
  - Useful when building sets of containers that cooperate





# VM vs containers



## Virtual Machines


- Slow boot/shutdown sequences
- Operating system overhead
- Total isolation between VMs
- Complete flexibility

## Containers

- Fast startup/teardown
- Small resource tracking overhead
- Fine-grained control of isolation
- Runs on host operating system

# Containers run on the host OS

- Applications within containers are limited to the host OS
- Cannot run Windows applications on Linux containers
  - Need a virtual machine for that
- Applications cannot use features not supported by the OS kernel
  - May require kernel updates





# CompSci 401: Cloud Computing

# Docker


Prof. Ítalo Cunha






# Docker containers

- Small containers running a single application
  - One container image per application
  - Small images speed up download
  - Single application speeds up startup time




# Docker containers

- Small containers running a single application
  - One container image per application
  - Small images speed up download
  - Single application speeds up startup time
- Libraries and dependencies bundled in image
  - No dependencies on host OS
  - Works on all OSes




# Docker containers

- Small containers running a single application
  - One container image per application
  - Small images speed up download
  - Single application speeds up startup time
- Libraries and dependencies bundled in image
  - No dependencies on host OS
  - Works on all OSes
- Read-only images
  - Changes to container to not persist upon restart
  - Reproducible execution
    - Container starts from a known state every time



# Building Docker containers

- Tooling around the development and deployment of containers
  - [Dockerfile](#) set of instructions
- Extensive image registry (DockerHub)
  - Ready-to-use containers/applications
  - Building blocks for more complex containers



# Dockerfile: a container recipe

Django uses Python

```
FROM python:3.7-slim-buster

# Install dependencies
RUN apt-get update
RUN apt-get install --no-install-recommends -y bird
RUN apt-get install --no-install-recommends -y libgit2-27 libgit2-dev
RUN apt-get install --no-install-recommends -y procps

# Copy application to container
RUN mkdir -p /usr/src/app
WORKDIR /usr/src/app
COPY . /usr/src/app

# Install modules in requirements.txt
RUN pip3 install --no-cache-dir -r requirements.txt

ENTRYPOINT entrypoint.sh
```

# Dockerfile: a container recipe

Django uses Python

We prefer Debian

```
FROM python:3.7-slim-buster

# Install dependencies
RUN apt-get update
RUN apt-get install --no-install-recommends -y bird
RUN apt-get install --no-install-recommends -y libgit2-27 libgit2-dev
RUN apt-get install --no-install-recommends -y procps

# Copy application to container
RUN mkdir -p /usr/src/app
WORKDIR /usr/src/app
COPY . /usr/src/app

# Install modules in requirements.txt
RUN pip3 install --no-cache-dir -r requirements.txt

ENTRYPOINT entrypoint.sh
```



# Dockerfile: a container recipe

Django uses Python

Install other applications and system components we need

We prefer Debian

```
FROM python:3.7-slim-buster

# Install dependencies
RUN apt-get update
RUN apt-get install --no-install-recommends -y bird
RUN apt-get install --no-install-recommends -y libgit2-27 libgit2-dev
RUN apt-get install --no-install-recommends -y procps

# Copy application to container
RUN mkdir -p /usr/src/app
WORKDIR /usr/src/app
COPY . /usr/src/app

# Install modules in requirements.txt
RUN pip3 install --no-cache-dir -r requirements.txt

ENTRYPOINT entrypoint.sh
```

# Dockerfile: a container recipe

We prefer Debian

Django uses Python

Install other applications and system components we need

Copy our application to the container

```
FROM python:3.7-slim-buster

# Install dependencies
RUN apt-get update
RUN apt-get install --no-install-recommends -y bird
RUN apt-get install --no-install-recommends -y libgit2-27 libgit2-dev
RUN apt-get install --no-install-recommends -y procps

# Copy application to container
RUN mkdir -p /usr/src/app
WORKDIR /usr/src/app
COPY . /usr/src/app

# Install modules in requirements.txt
RUN pip3 install --no-cache-dir -r requirements.txt

ENTRYPOINT entrypoint.sh
```

# Dockerfile: a container recipe

We prefer Debian

Django uses Python

```
FROM python:3.7-slim-buster
```

Install other applications and system components we need

```
# Install dependencies
```

```
RUN apt-get update
```

```
RUN apt-get install --no-install-recommends -y bird
```

```
RUN apt-get install --no-install-recommends -y libgit2-27 libgit2-dev
```

```
RUN apt-get install --no-install-recommends -y procs
```

Copy our application to the container

```
# Copy application to container
```

```
RUN mkdir -p /usr/src/app
```

```
WORKDIR /usr/src/app
```

```
COPY . /usr/src/app
```

Install Python modules (including Django)

```
# Install modules in requirements.txt
```

```
RUN pip3 install --no-cache-dir -r requirements.txt
```

```
ENTRYPOINT entrypoint.sh
```

# Dockerfile: a container recipe

We prefer Debian

Django uses Python

```
FROM python:3.7-slim-buster
```

Install other applications and system components we need

```
# Install dependencies
```

```
RUN apt-get update
```

```
RUN apt-get install --no-install-recommends -y bird
```

```
RUN apt-get install --no-install-recommends -y libgit2-27 libgit2-dev
```

```
RUN apt-get install --no-install-recommends -y procps
```

Copy our application to the container

```
# Copy application to container
```

```
RUN mkdir -p /usr/src/app
```

```
WORKDIR /usr/src/app
```

```
COPY . /usr/src/app
```

Install Python modules (including Django)

```
# Install modules in requirements.txt
```


```
RUN pip3 install --no-cache-dir -r requirements.txt
```

Set the application to run inside the container

```
ENTRYPOINT entrypoint.sh
```

# Docker container runtime

- Container runtime controls container
- User issues commands to the runtime using CLI tools or network APIs
- Runtime sets up required components in the operating system
  - Creates and destroys namespaces
  - Configures namespaces
    - For example, the network and volumes



# Open source and standards

- Docker is a toolset
- The container runtime has been standardized and open-sourced
  - containerd.io
- Other controllers for the container runtime exist
  - Kubernetes
  - Red Hat Podman



# Linux Containers (LXC)

- LXC is another tool to run containers on Linux
  - Relies on the same technologies (namespaces)
- Different approach from Docker
  - More similar to a virtual machine
  - Larger images with a complete user environment
  - Multiple applications running in the container
  - State persists between executions of the same container