



Monitoring and Telemetry

Prof. Ítalo Cunha

Telemetry is fundamental in cloud operations

- Monitor system behavior
- Evaluate performance
- Forecast demand
- Provision resources
- Troubleshoot failures
- Identify attacks
- Diagnose problems
- Perform audits
- Accurate billing

Metrics

- **Quantitative** data about a system
 - Link bandwidth, CPU utilization, memory usage
 - May be categorical: status (up/down), properties
- Collected periodically
 - Active: Run a task to collect a measurement (ping for RTT)
 - Passive: Read or receive data (CPU utilization or SNMP with link utilization)
- Metrics can be associated with
 - Physical resources: servers, switches, links
 - Virtual resources: VMs, containers, virtual disks, virtual networks
 - Abstractions: Deployment, load balancer, databases

Logs

- **Qualitative** data about noteworthy events
 - Can be used proactively to identify issues and trigger alerts
 - More often used reactively to troubleshoot issues after detection
- Logging messages should have some standard to ease processing
 - Timestamp
 - Identification of the code generating the message

Logs

- **Qualitative** data about noteworthy events
 - Can be used proactively to identify issues and trigger alerts
 - More often used reactively to troubleshoot issues after detection
- Logging messages should have some standard to ease processing
 - Timestamp
 - Identification of the code generating the message

Linux dmesg

```
[ 0.473042] thermal_sys: Registered thermal governor 'user_space'
[ 0.473042] thermal_sys: Registered thermal governor 'power_allocator'
[ 0.473042] cpuidle: using governor ladder
[ 0.473042] cpuidle: using governor menu
[ 0.473042] ACPI FADT declares the system doesn't support PCIe ASPM, so disable it
[ 0.473042] ACPI: bus type PCI registered
[ 0.473042] acpiphp: ACPI Hot Plug PCI Controller Driver version: 0.5
[ 0.473042] PCI: MMCONFIG for domain 0000 [bus 00-ff] at [mem 0xe0000000-0xffffffff]
[ 0.473042] PCI: MMCONFIG at [mem 0xe0000000-0xffffffff] reserved in E820
```


Logs

- **Qualitative** data about noteworthy events
 - Can be used proactively to identify issues and trigger alerts
 - More often used reactively to troubleshoot issues after detection
- Logging messages should have some standard to ease processing
 - Timestamp
 - Identification of the code generating the message

Linux dmesg

```
[ 0.473042] thermal_sys: Registered thermal governor 'user_space'
[ 0.473042] thermal_sys: Registered thermal governor 'power_allocator'
[ 0.473042] cpuidle: using governor ladder
[ 0.473042] cpuidle: using governor menu
[ 0.473042] ACPI FADT declares the system doesn't support PCIe ASPM, so disable it
[ 0.473042] ACPI: bus type PCI registered
[ 0.473042] acpi: ACPI Hot Plug GPF Controller Driver Version 0.5
[ 0.473042] 127.0.0.1 user-identifier frank [10/Oct/2000:13:55:36 -0700] "GET /apache_pb.gif HTTP/1.0" 200 2326
[ 0.473042] PCI: MMCONFIG at [mem 0xe0000000-0xffffffff] reserved in E820
```

Common log
format

Traces

- Records of relationships during the execution of a task
 - Logs capture events
 - Traces capture the sequence and relationships between events or steps
 - Fine grained information
- Traces are collected across process boundaries
 - A distributed challenge
 - Interactions may not be known in advance
 - Synchronization across devices is important for troubleshooting time-dependent issues

Traces

- Records of relationships during the execution of a task
 - Logs capture events
 - Traces capture the sequence and relationships between events or steps
 - Fine grained information
- Traces are collected across process boundaries
 - A distributed challenge
 - Interactions may not be known in advance
 - Synchronization across dependent issues

```
zeus:~ $ strace ls
execve("/bin/ls", ["ls"], 0x7ffe5b1b0cc0 /* 48 vars */) = 0
brk(NULL)                                = 0x5629686f5000
access("/etc/ld.so.preload", R_OK)       = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=97703, ...}) = 0
mmap(NULL, 97703, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7ffb7d6ff000
close(3)                                 = 0
openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libselinux.so.1", O_RDONLY|O_CLOEXEC) = 3
```


Telemetry use cases

Monitoring

- Ongoing measurements
- Proactive
- Needs to be lightweight
- Needs broad coverage
- Needs to be automated

Troubleshooting

- On-demand measurements
- Reactive
- May be expensive
- May be specific
- May require human oversight

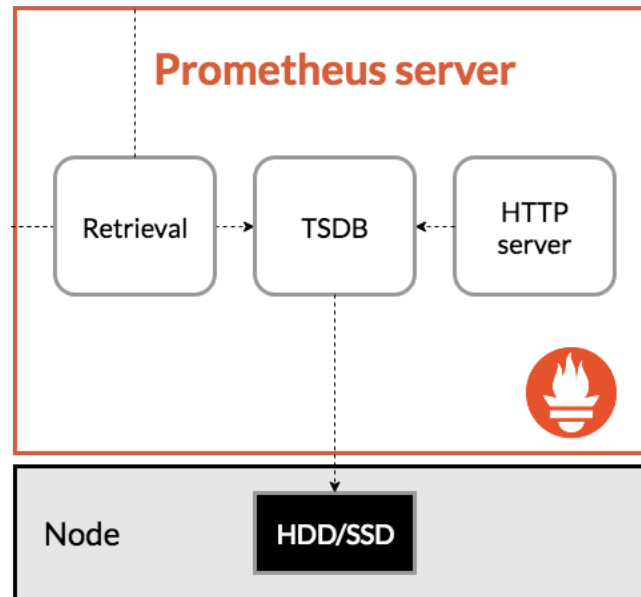
Monitoring needs tooling to be useful

- Automated collection, processing, indexing, and storage
- Alerts to detect events or conditions
- Closed-loop controllers that can take actions depending on conditions
- Dashboards for human consumption

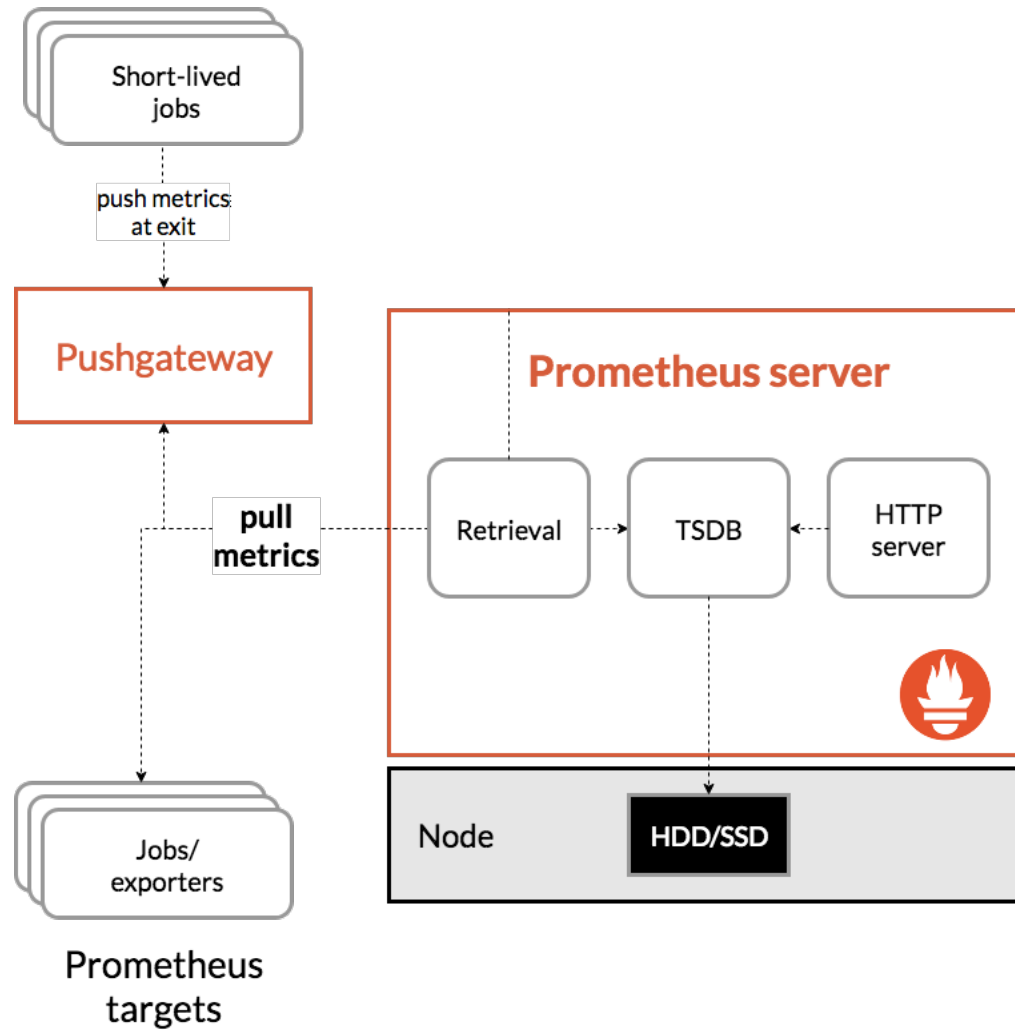
Monitoring needs tooling to be useful

- Automated collection, processing, indexing, and storage
 - Alerts to detect events or conditions
 - Closed-loop controllers that can take actions depending on conditions
 - Dashboards for human consumption
-
- Integration with machine learning for anomaly detection and inference

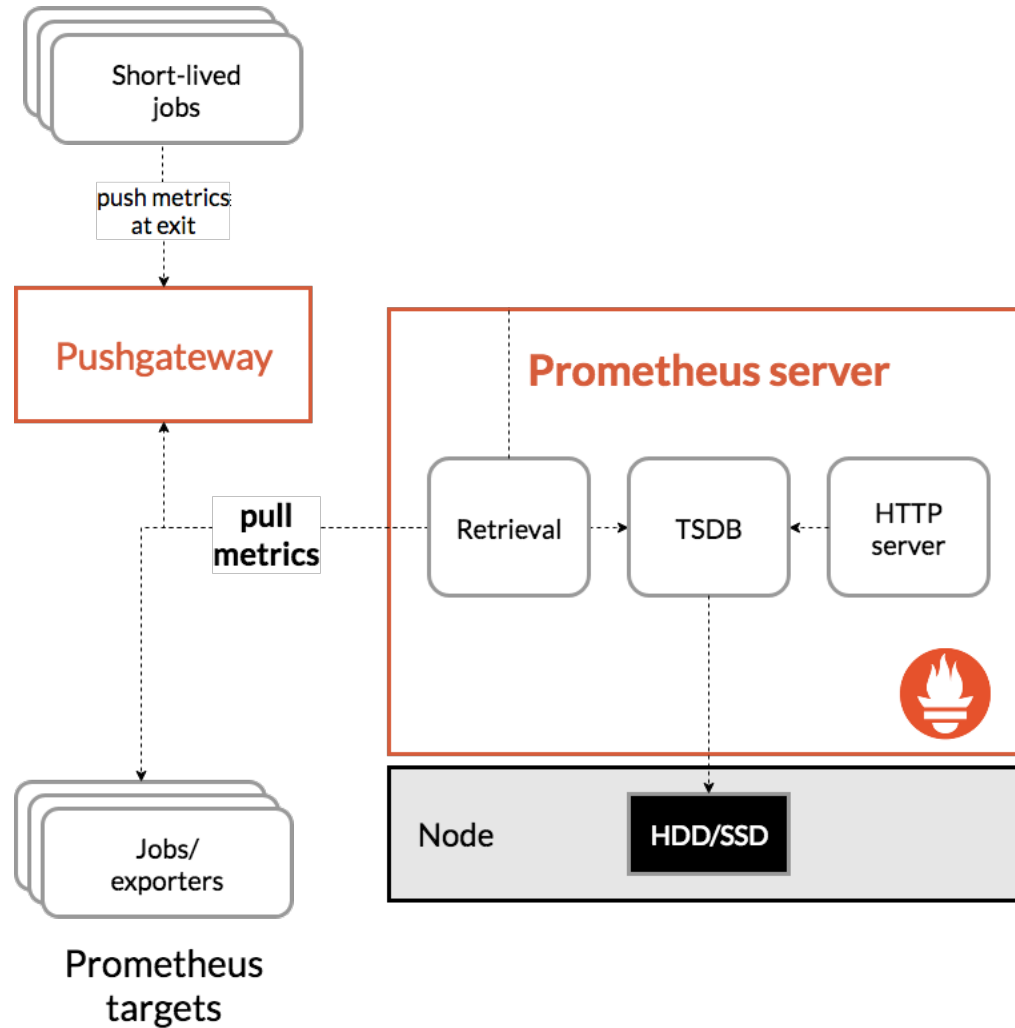
Prometheus



Prometheus

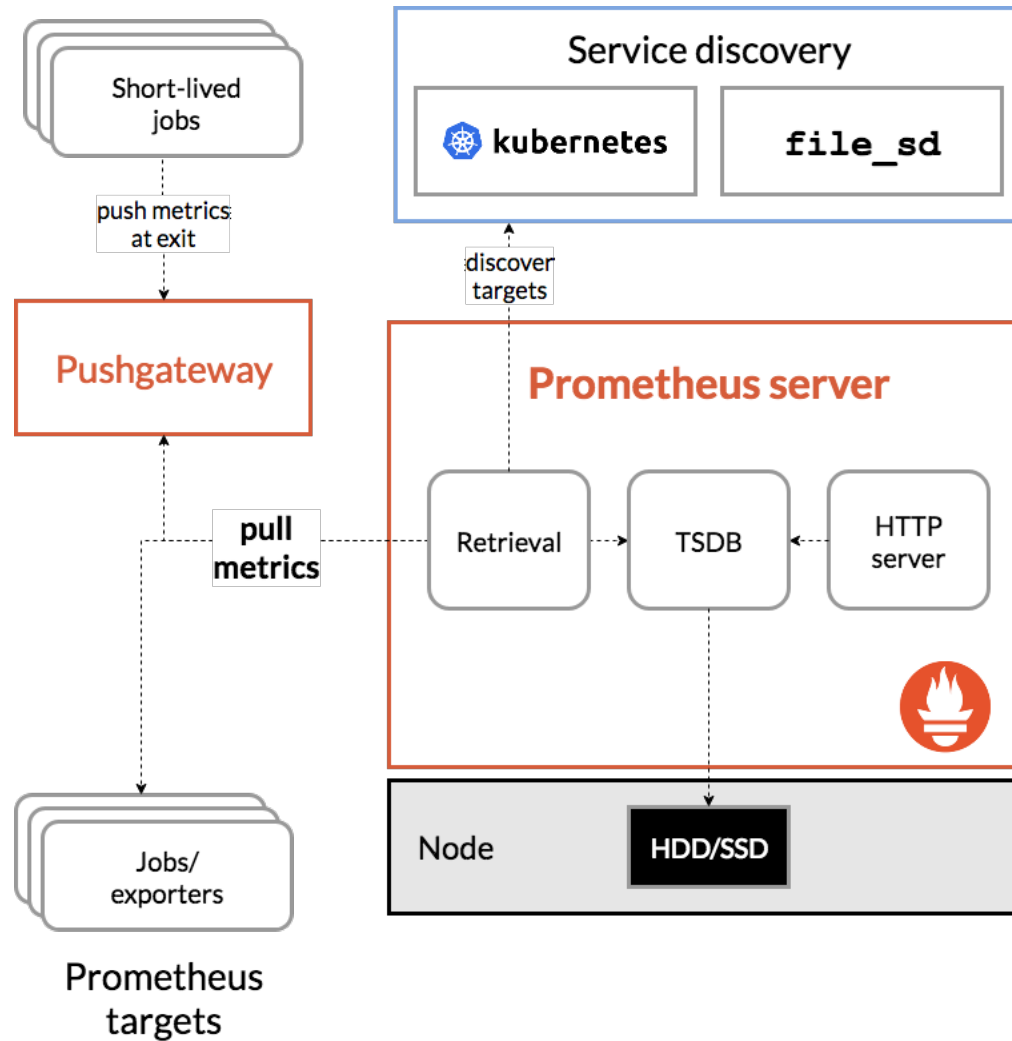


Prometheus

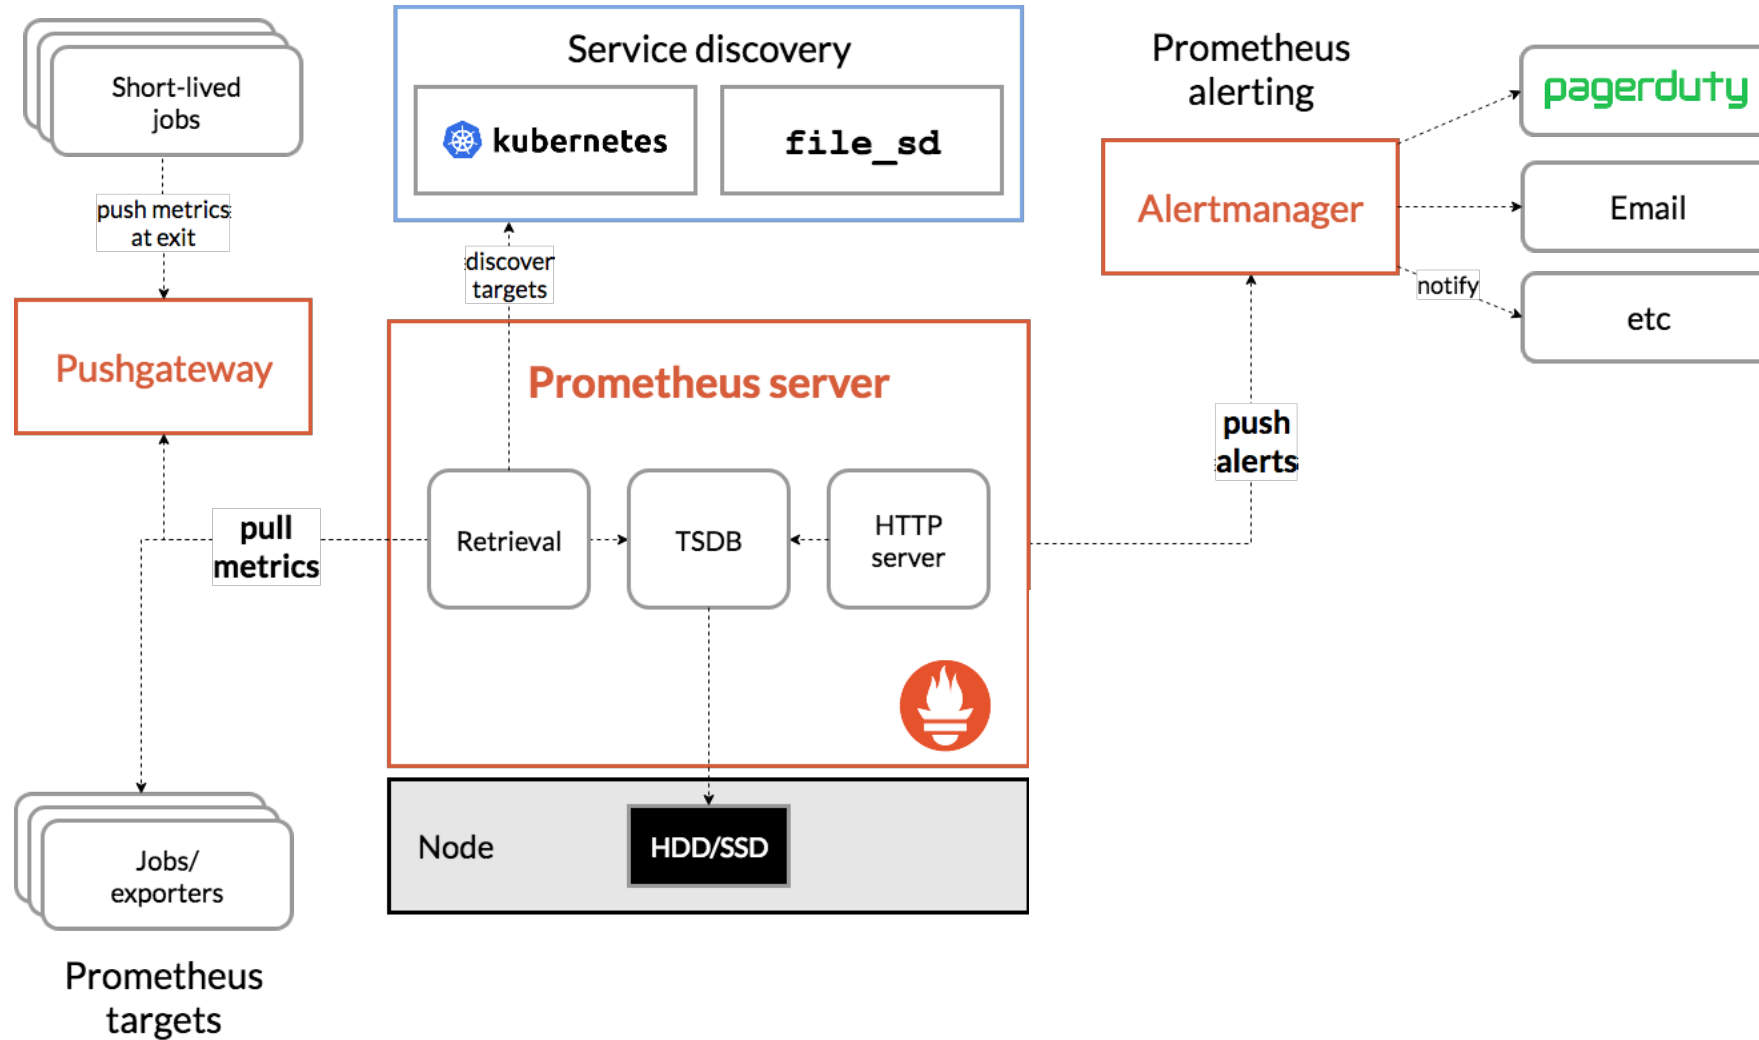


Client
libraries

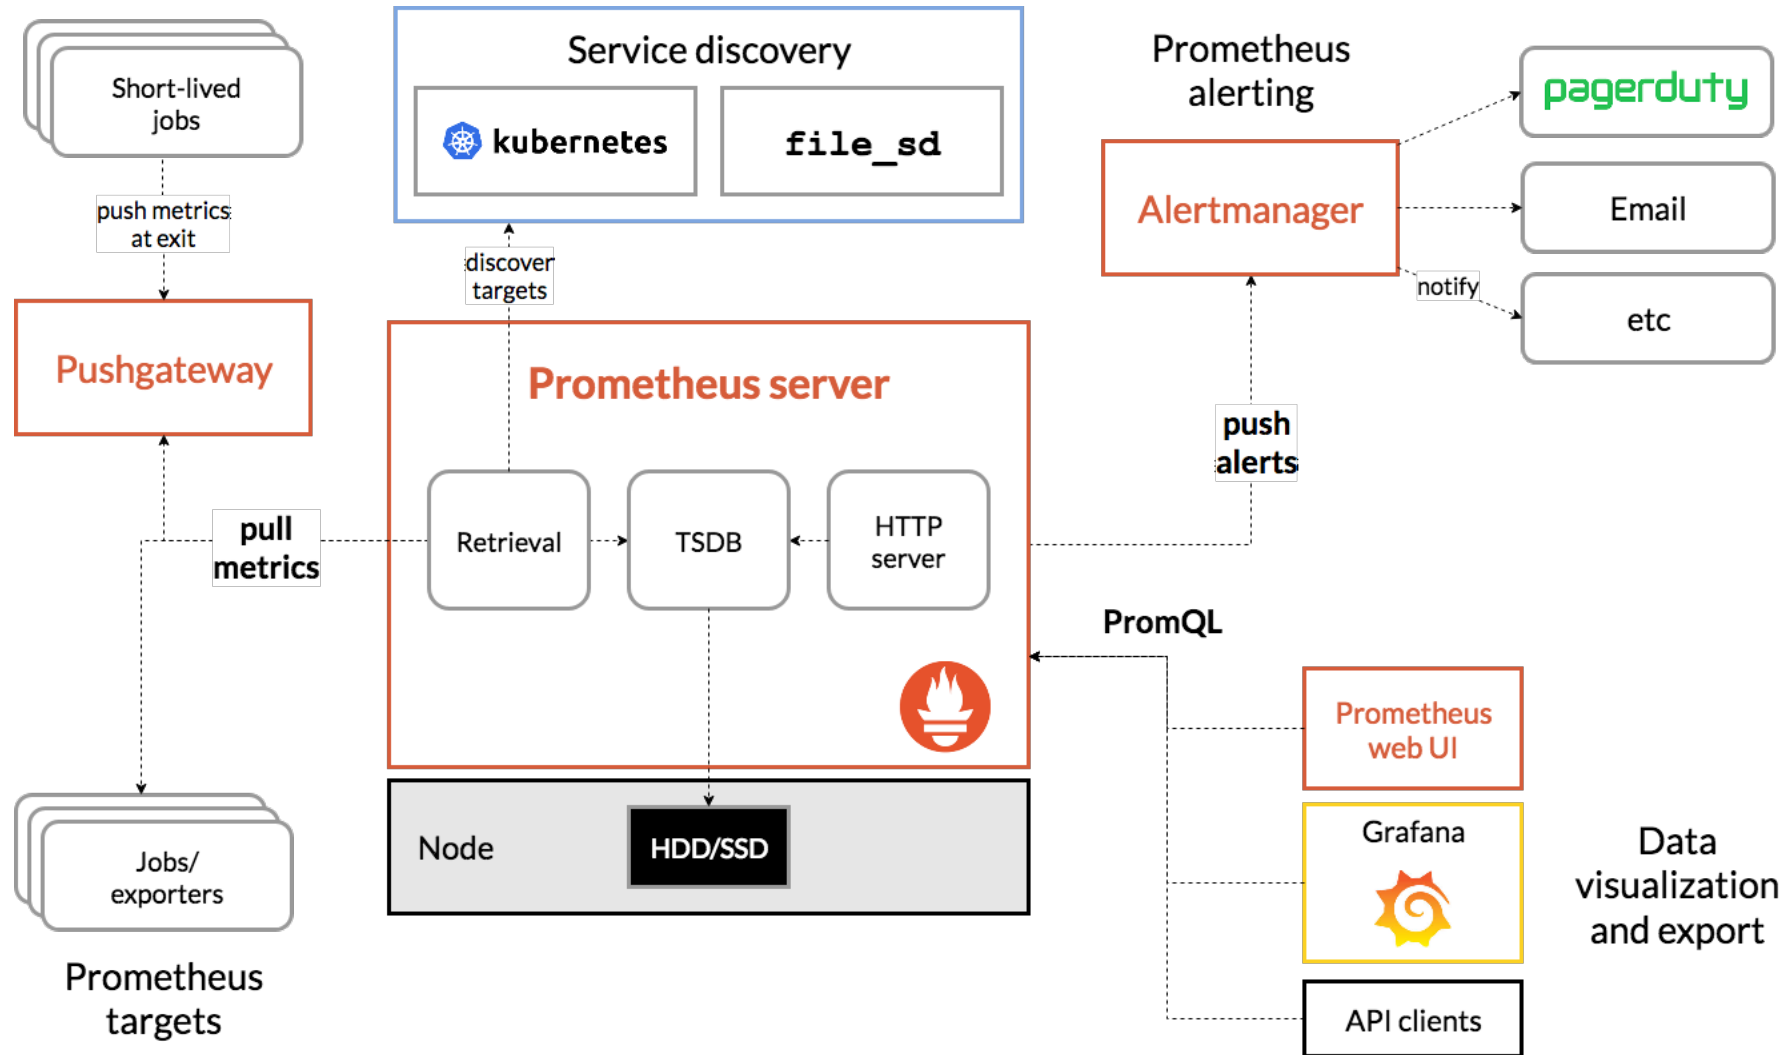
Prometheus



Prometheus



Prometheus



Handling distributed infrastructures

- Where to store monitoring and telemetry data?
 - Locally on each site vs centralized on a single site
- Trade-offs to consider
 - Possible use of bandwidth to transfer large amounts of data
 - Replication requires storage space
 - Accessing remote data may add delays

Handling distributed infrastructures

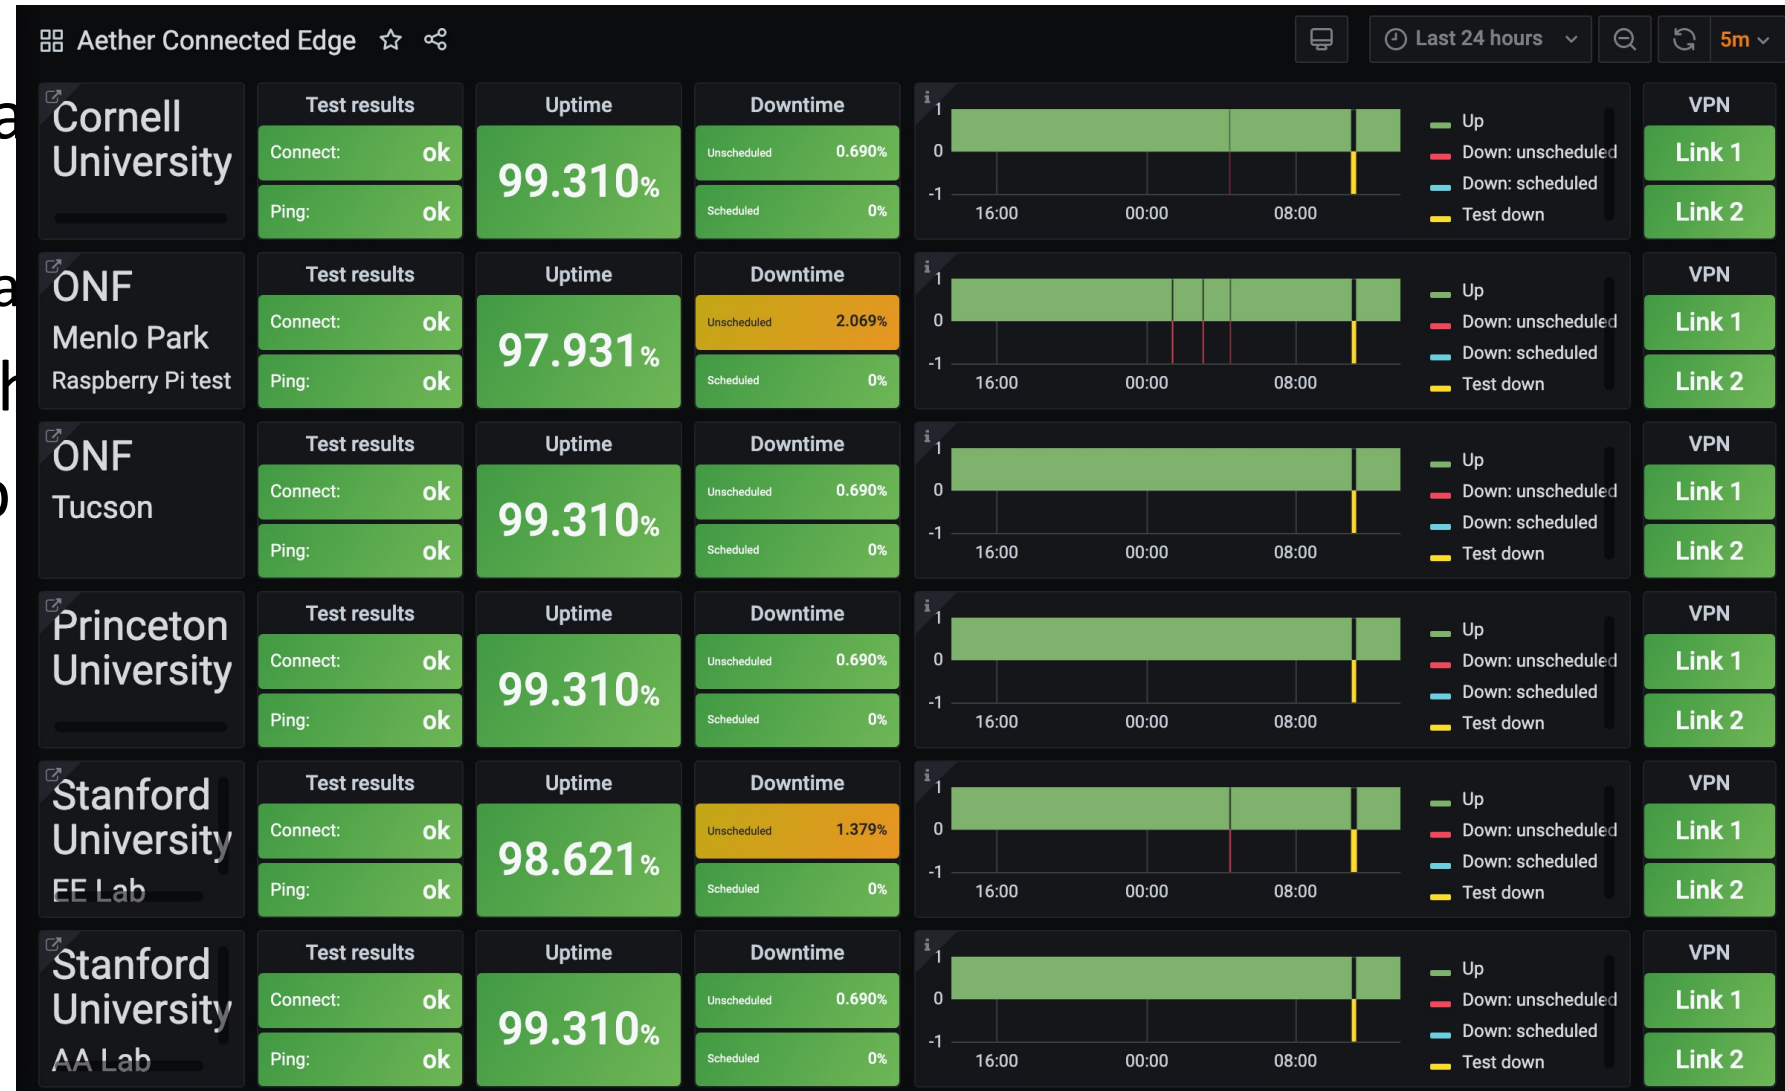
- Where to store monitoring and telemetry data?
 - Locally on each site vs centralized on a single site
- Trade-offs to consider
 - Possible use of bandwidth to transfer large amounts of data
 - Replication requires storage space
 - Accessing remote data may add delays
- Aether uses a mixes approach
 - Keeps raw data local but centralizes query results

Graphana

- Predefined visualizations for different types of data
 - Time series
 - Resource allocations and utilizations
- Integrations with telemetry frameworks and applications
- A Graphana deployment can pull information from multiple sources


Graphana

- Predefined visualizations
 - Time series
 - Resource allocation
- Integrations with
- A Graphana deployment



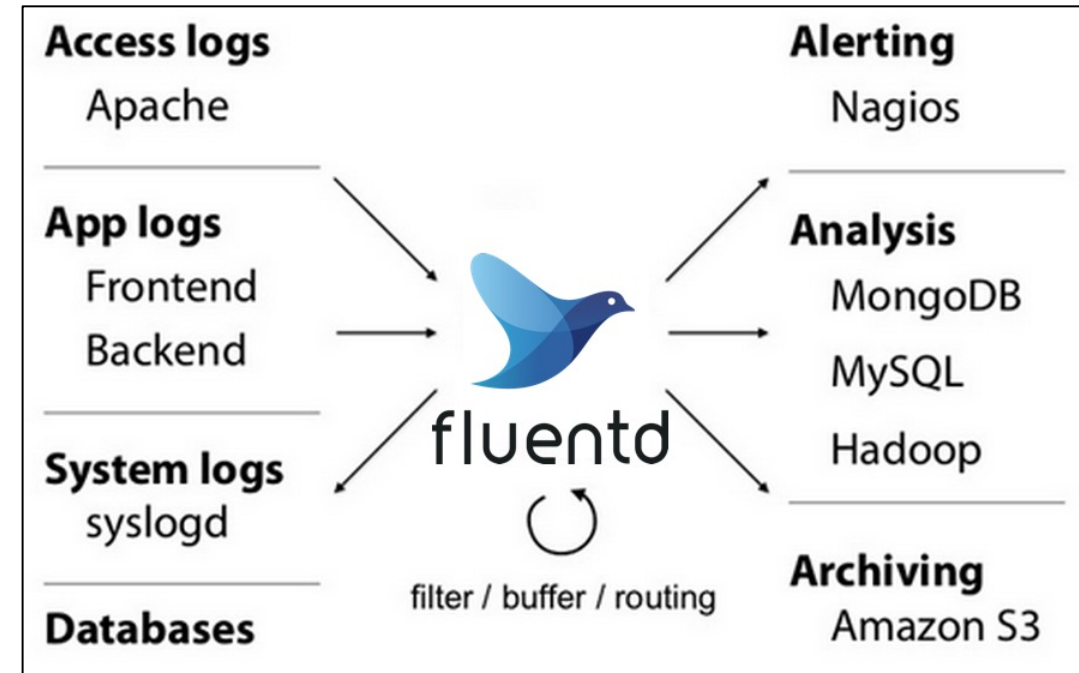

Logging for cloud-native applications

- Need to aggregate and correlate system and application logs



Logging for cloud-native applications

- Need to aggregate and correlate system and application logs



Logging for cloud-native applications

- Need to aggregate and correlate system and application logs
- Need to normalize logs

```
2020-08-18 05:35:54.842Z INFO [DistributedP4RuntimeTableMirror]
Synchronized TABLE_ENTRY mirror for device:leaf1: 0 removed, 2 updated, 4 added
```



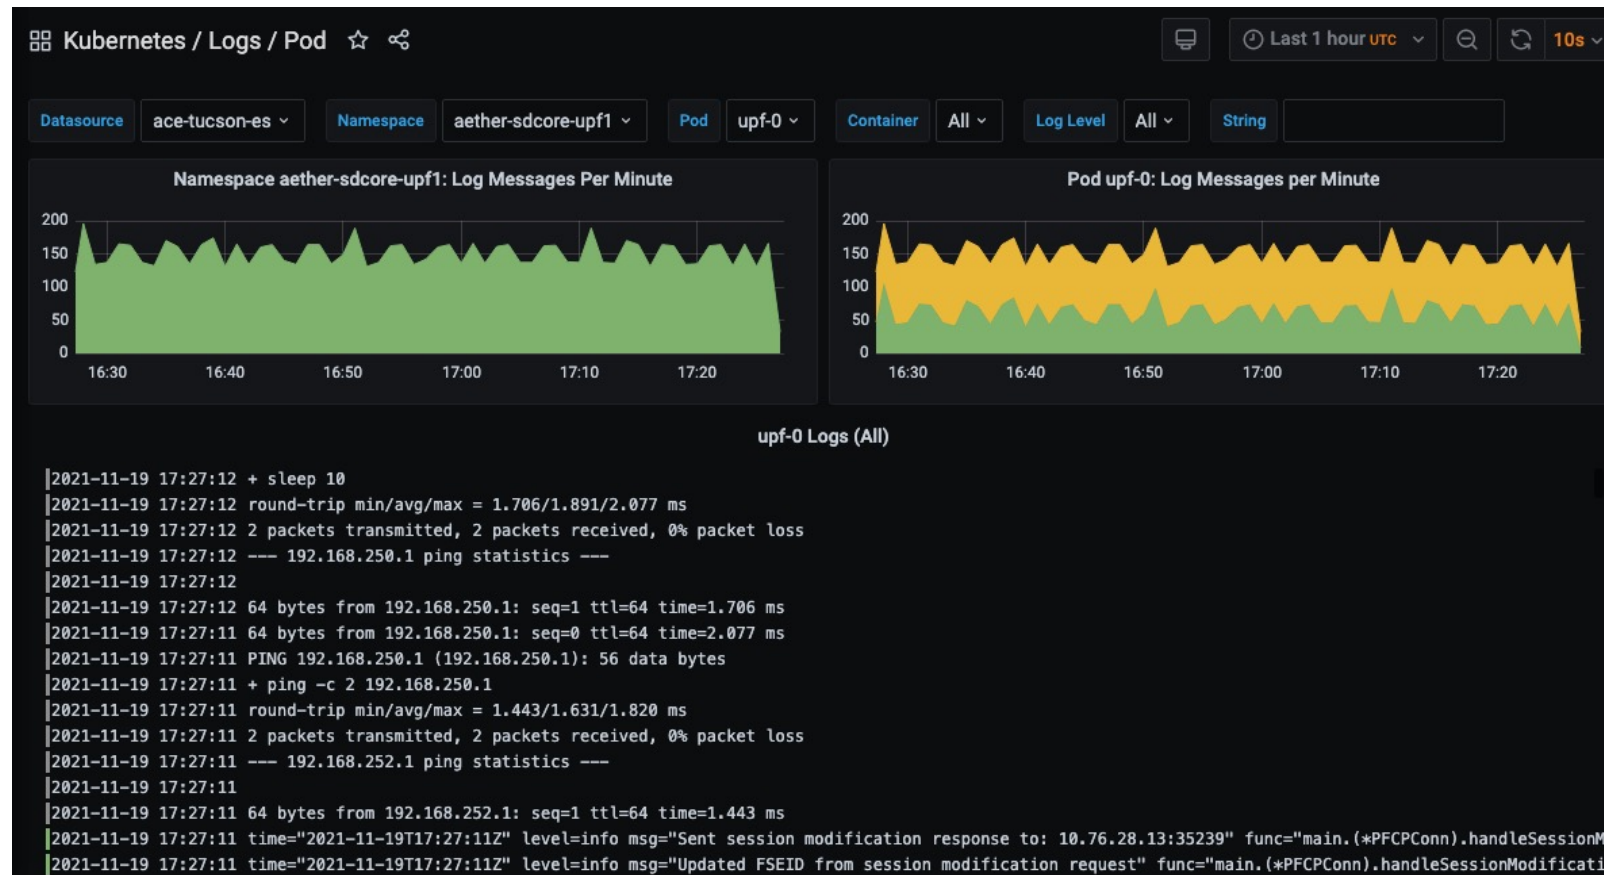
```
{
  "time": "2020-08-18 05:35:54.842Z",
  "logLevel": "INFO",  "component": "DistributedP4RuntimeTableMirror",
  "log": "Synchronized TABLE_ENTRY mirror for device:leaf1: 0 removed, 2 updated, 4 added"
}
```

Best practices

- Log shipping handled by the platform
 - Use if a logging library or stdout/stderr
- File logging should be disabled
 - Layered filesystems used in containers have poor write performance
 - Unnecessary if logging is handled by the platform
- Asynchronous logging is preferred
- Timestamps should be accurate
 - Stamp as close as possible to log event, keeps clocks synchronized
- Ability to change log levels at runtime
- Integrated dashboards

Integrated dashboards

- Combine monitoring, logging, and tracing in a single interface



Tracing an application

- Need to collect fine-grained information from processes
- And keep track of requests across process boundaries
- Tracing frameworks integrate with existing code
 - Language runtimes (e.g., JVM, node.js)
 - Frameworks (e.g., Flask, Spring, ASP.NET)

Observability

- New buzzword for monitoring and telemetry
- Denotes a system that follows the best practices we have discussed
 - Many recent innovations in this area, systems need to integrate
 - Software frameworks for host-based monitoring
 - Service meshes
 - Software-defined networks with programmable, on-path monitoring

